

Hashing: SUHA: $s \Theta(1 + \alpha)$ — the 1 comes from applying the hash function and random access to the slot whereas the α comes from searching the list. This is equal to $O(1)$ if $\alpha = O(1)$, i.e., $m = \Omega(n)$.

Chaining: $\alpha = n/m$. We denote α as the load factor of the hash table. By ensuring m is large enough to keep α small, searches and deletions in a hash table using chaining take $O(1 + \alpha)$ or approximately $O(1)$ time.

Open addressing: alternative location in the hash table, uniform hashing assumption says that "Each key k is equally likely to hash to any of the m ! permutations of $\langle 0, 1, \dots, (m - 1) \rangle$ ", **linear probing** tries consecutive locations one by one until it finds a free one, The expected number of lookups is $1/(1 - \alpha)$, which can become very large when α is close to 1.

Q: Consider an **open addressing** hash table with m slots, where collisions are handled using **linear probing**. Assume simple uniform hashing, and assume that the table is initially empty. The probability that the first two slots of the table are filled after the first two insertions is $3/m^2$. TRUE There are m^2 equally likely ways to map the first two keys into the table. Of those, the following three possibilities lead to the first two slots being filled: $h(K1) = 0, h(K2) = 1$; $h(K1) = 1, h(K2) = 0$; and $h(K1) = 0, h(K2) = 0$.

Double hashing: $hm = (h1(k) + i \cdot h2(k)) \bmod m$

Rolling hash: Rabin-Karp: $\text{hash}() \rightarrow$ computes the hash of the list. • $\text{append}(\text{val}) \rightarrow$ adds val to the end of the list. • $\text{skip}(\text{val}) \rightarrow$ removes the front element from the list, assuming it is val .

Graph/ Search Material on next page.

Bellman-Ford Analysis

```

for v in V:
    v.d = ∞
    v.π = None
s.d = 0
for i from 1 to |V| - 1:
    for (u, v) in E:
        relax(u, v)
for (u, v) in E:
    if v.d > u.d + w(u, v):
        report that a negative-weight cycle exists
    
```

Complexity Analysis:

- Initialization: $O(V)$
- Relaxation: $O(E)$ per iteration, $O(V)$ iterations → $O(VE)$
- Total: $O(VE)$

DFS: not sorting, run time: $O(E + V)$, space: $O(V)$, **stack**, maze, for DAGs, non-neg edges. **Not for shortest detection**, does: detect cycles (no back edges, no cycles), topological sort. DFS tree, $V-1$ edges in DFS tree. No cross-edges in dfs tree for undirected

Triangle inequality: For any edge (u, v) , we have $\delta(s, v) \leq \delta(s, u) + w(u, v)$. In English, the weight of the shortest path from s to v is no greater than the weight of the shortest path from s to u plus the weight of the edge from u to v

DAG-SP: sorting DAGs, $O(V + E)$, We know that there will be no path from a vertex v to another vertex u occurring before v in the topological order. The first step in finding the shortest path, therefore, should be to topologically sort the graph, in $O(V + E)$ time. Now we go through each vertex starting from the beginning, and relax each outgoing edge. This will guarantee that each edge only need be relaxed once, making the total runtime $O(V + E)$. A sample execution is shown in Figure 1.

• **Bellman Ford: sorting, Negative weight edges:** For graphs that might have negative weight edges we use Bellman Ford's algorithm to find all shortest paths from a source vertex. Bellman Ford relaxes all $O(|E|)$ edges of a graph exactly $O(|V|)$ times, so the running time of the algorithm is $O(|V| |E|)$. Bellman Ford can also detect if a graph has a negative weight cycle reachable from the source vertex or not, and if so, does not return a shortest path.

```

DIJKSTRA (G, W, s) //uses priority queue Q
Initialize (G, s)
S ← ∅
Q ← V[G] //Insert into Q
while Q ≠ ∅
    do u ← EXTRACT-MIN(Q) //deletes u from Q
    S = S ∪ {u}
    for each vertex v ∈ Adj[u]
        do RELAX (u, v, w) ← this is an implicit DECREASE_KEY operation
    
```

BFS: sort, unweighted $O(V + E)$, priority queue, undirected, thru 'levels', going from neighbors of s , to neighbors or neighbors, etc. BFS tree, $V-1$ edges in BFS tree

- $v.d$ - The weight of the current shortest path from s to v .
- $v.\pi$ - The parent vertex of v in the current shortest path.
- $w(u, v)$ is the weight of the edge from vertex u to vertex v
- $\delta(u, v)$ is the weight of the shortest path from vertex u to vertex v

Dijkstra: sorting weighted, non-neg, $O(V \log V + E)$, weighted- directed graphs, only guaranteed for non-negative weight edges. min-heap fibonacci (priority queue), calc from min & remove & put in 'visited' until, reach goal state/vertex. sorting, runtime: $O(B + X|V| + D|E|)$, where B is the time to build the priority queue, X is the time for *EXTRACT-MIN* and D is the *DECREASE-KEY* operation, cant fix dijkstra with making weights positive, Dijkstra's algorithm will return an incorrect answer only when a negative edge goes into an already processed node and there are edges coming out of that vertex

Fibonacci-heap: insert = $\theta(1)$, dec-key = $\theta(1)$, extract-min = $\theta(\log n)$